

## Practical 1

---

**Aim: Hadoop Configuration and Single node cluster setup and perform file management task in Hadoop**

**Aim: Hadoop Configuration and Single node cluster setup and perform file management task in**

**Hadoop. (creating a directory, list the content of directory, upload and download file in HDFS)**

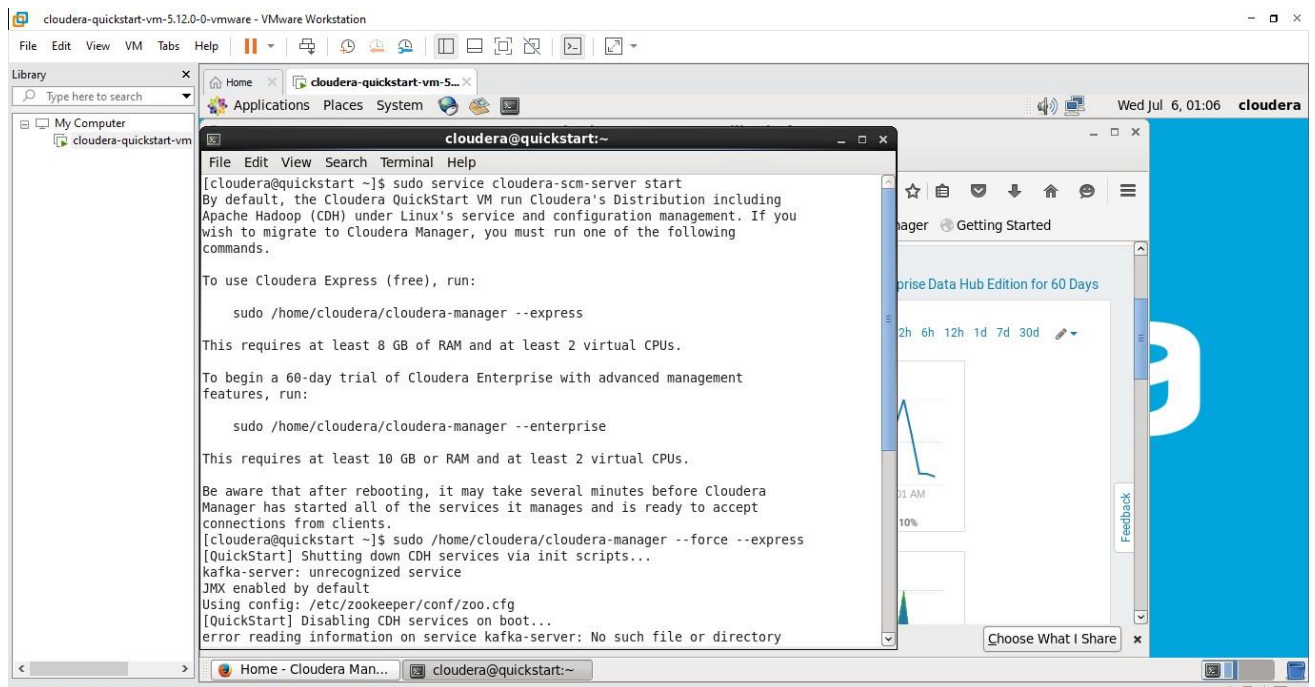
- Prerequisite: dos/linux command
- Configure ubuntu/cloudera
- Test Hadoop services

**172.16.12.38 - /**

---

7/5/2022	1:04 PM	6578981376	<a href="#">cloudera-quickstart-vm-5.12.0-0-vmware.ova</a>
7/15/2021	4:36 PM	5811085017	<a href="#">cloudera-quickstart-vm-5.12.0-0-vmware.zip</a>
9/26/2019	1:07 PM	536135688	<a href="#">VMware-workstation-full-15.0.4-12990004.exe</a>
9/28/2019	1:26 PM	434232619	<a href="#">VMware_Workstation_14_Pro.zip</a>
7/5/2022	12:59 PM	301	<a href="#">web.config</a>

---



```

cloudera@quickstart:~
File Edit View Search Terminal Help
[cloudera@quickstart ~]$ sudo service cloudera-scm-server start
By default, the Cloudera QuickStart VM run Cloudera's Distribution including
Apache Hadoop (CDH) under Linux's service and configuration management. If you
wish to migrate to Cloudera Manager, you must run one of the following
commands.

To use Cloudera Express (free), run:

    sudo /home/cloudera/cloudera-manager --express

This requires at least 8 GB of RAM and at least 2 virtual CPUs.

To begin a 60-day trial of Cloudera Enterprise with advanced management
features, run:

    sudo /home/cloudera/cloudera-manager --enterprise

This requires at least 10 GB of RAM and at least 2 virtual CPUs.

Be aware that after rebooting, it may take several minutes before Cloudera
Manager has started all of the services it manages and is ready to accept
connections from clients.
[cloudera@quickstart ~]$ sudo /home/cloudera/cloudera-manager --force --express
[QuickStart] Shutting down CDH services via init scripts...
kafka-server: unrecognized service
JMX enabled by default
Using config: /etc/zookeeper/conf/zoo.cfg
[QuickStart] Disabling CDH services on boot...
error reading information on service kafka-server: No such file or directory

```

```

cloudera@quickstart:~
File Edit View Search Terminal Help
Manager has started all of the services it manages and is ready to accept
connections from clients.
[cloudera@quickstart ~]$ sudo /home/cloudera/cloudera-manager --force --express
[QuickStart] Shutting down CDH services via init scripts...
kafka-server: unrecognized service
JMX enabled by default
Using config: /etc/zookeeper/conf/zoo.cfg
[QuickStart] Disabling CDH services on boot...
error reading information on service kafka-server: No such file or directory
[QuickStart] Starting Cloudera Manager server...
[QuickStart] Waiting for Cloudera Manager API...
[QuickStart] Starting Cloudera Manager agent...
[QuickStart] Configuring deployment...
Submitted jobs: 15
[QuickStart] Deploying client configuration...
Submitted jobs: 16
[QuickStart] Starting Cloudera Management Service...
Submitted jobs: 24
[QuickStart] Enabling Cloudera Manager daemons on boot...

Success! You can now log into Cloudera Manager from the QuickStart VM's browser:

http://quickstart.cloudera:7180

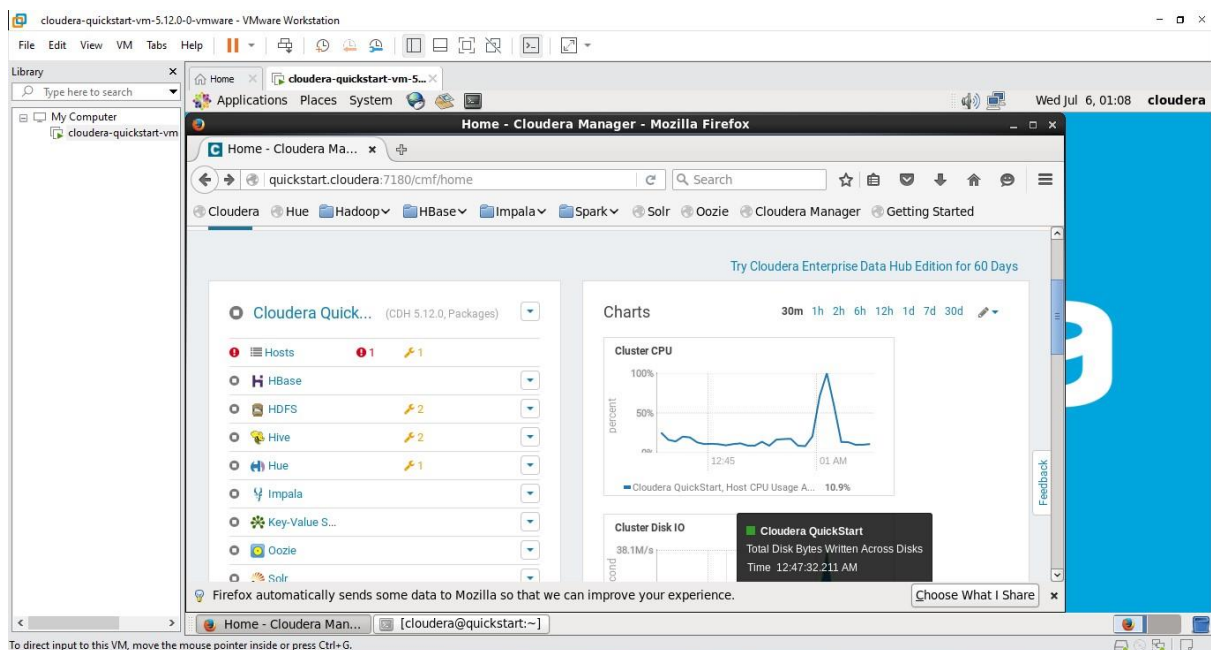
Username: cloudera
Password: cloudera

[cloudera@quickstart ~]$

```

Login user and password

And check dashboard



The screenshot shows a web interface with a list of services. Each service has a radio button, an icon, a name, a status indicator (red circle with 'i' or yellow key), a count, and a dropdown arrow. The services listed are: Hosts, HBase, HDFS, Hive, Hue, Impala, Key-Value S..., Oozie, Solr, Spark, Sqoop 1 Cl..., Sqoop 2, and YARN (MR2...). Below the list, a grey bar contains the text: "Firefox automatically sends so".

Service	Status	Count
Hosts	1	1
HBase		
HDFS	2	
Hive	2	
Hue	1	
Impala		
Key-Value S...		
Oozie		
Solr		
Spark		
Sqoop 1 Cl...		
Sqoop 2		
YARN (MR2...		

Firefox automatically sends so

## Practical 2

---

**Aim: Copy your data into the Hadoop Distributed File System (HDFS), creating a directory, list the content of directory, upload and download file in HDFS**

- Prerequisite: dos/linux command
- Configure HDFS

Renaming a file in hdfs:

```
[cloudera@quickstart ~]$ sudo -u hdfs hadoop fs -mv /user/19ce063/arpan.txt /user/19ce063/19ce063_arpan.txt
[cloudera@quickstart ~]$ hadoop fs -ls /user/19ce063/
Found 1 items
-rw-r--r-- 1 hdfs supergroup      18 2022-07-28 00:00 /user/19ce063/19ce063_arpan.txt
```

Copying file in hdfs:

```
[cloudera@quickstart ~]$ sudo -u hdfs hadoop fs -cp /user/19ce063/19ce063_arpan.txt /user/19ce063/copy-arpan.txt
22/07/28 00:07:42 WARN hdfs.DFSClient: Caught exception
java.lang.InterruptedException
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1281)
    at java.lang.Thread.join(Thread.java:1355)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFSOutputStream.java:952)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endBlock(DFSOutputStream.java:690)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStream.java:879)
[cloudera@quickstart ~]$ hadoop fs -ls /user/19ce063/
Found 2 items
-rw-r--r-- 1 hdfs supergroup      18 2022-07-28 00:00 /user/19ce063/19ce063_arpan.txt
-rw-r--r-- 1 hdfs supergroup      18 2022-07-28 00:07 /user/19ce063/copy-arpan.txt
```

Copying file from another folder:



```
[cloudera@quickstart ~]$ sudo -u hdfs hadoop fs -cp /user/19ce063/19ce063_arpan.txt /user/copying_in_anotherfolder.txt
[cloudera@quickstart ~]$ hadoop fs -ls /user
Found 12 items
drwxr-xr-x - 19ce063 supergroup          0 2022-07-28 00:07 /user/19ce063
drwxr-xr-x - cloudera cloudera          0 2022-07-21 00:50 /user/cloudera
-rw-r--r-- 1 hdfs supergroup          18 2022-07-28 00:10 /user/copying_in_anotherfolder.txt
drwx----- - hdfs supergroup          0 2022-07-21 00:57 /user/hdfs
drwxr-xr-x - hdfs supergroup          0 2022-07-21 00:32 /user/hduser
drwxr-xr-x - mapred hadoop            0 2017-07-19 06:29 /user/history
drwxrwxrwx - hive supergroup          0 2017-07-19 06:31 /user/hive
drwxrwxrwx - hue supergroup           0 2017-07-19 06:30 /user/hue
drwxrwxrwx - jenkins supergroup        0 2017-07-19 06:29 /user/jenkins
drwxrwxrwx - oozie supergroup          0 2017-07-19 06:30 /user/oozie
drwxrwxrwx - root supergroup           0 2017-07-19 06:29 /user/root
drwxr-xr-x - hdfs supergroup           0 2017-07-19 06:31 /user/spark
```

## Copying folder into another folder:

```
[cloudera@quickstart ~]$ sudo -u hdfs hadoop fs -mkdir /user/dummy
[cloudera@quickstart ~]$ sudo -u hdfs hadoop fs -mv /user/dummy /user/19ce063
[cloudera@quickstart ~]$ hadoop fs -ls /user/19ce063
Found 3 items
-rw-r--r-- 1 hdfs supergroup          18 2022-07-28 00:00 /user/19ce063/19ce063_arpan.txt
-rw-r--r-- 1 hdfs supergroup          18 2022-07-28 00:07 /user/19ce063/copy-arpan.txt
drwxr-xr-x - hdfs supergroup          0 2022-07-28 00:12 /user/19ce063/dummy
```

## Reading copied content file:

```
[cloudera@quickstart ~]$ hadoop fs -cat /user/19ce063/copy-arpan.txt
I am Arpan Ladani
[cloudera@quickstart ~]$ █
```

## To delete trash:

```
[cloudera@quickstart ~]$ sudo -u hdfs hadoop fs -expunge
22/07/28 00:33:54 INFO fs.TrashPolicyDefault: TrashPolicyDefault#deleteCheckpoint for trashRoot: hdfs://quickstart.cloudera:8020/user/hdfs/.Trash
22/07/28 00:33:54 INFO fs.TrashPolicyDefault: TrashPolicyDefault#deleteCheckpoint for trashRoot: hdfs://quickstart.cloudera:8020/user/hdfs/.Trash
22/07/28 00:33:54 INFO fs.TrashPolicyDefault: TrashPolicyDefault#createCheckpoint for trashRoot: hdfs://quickstart.cloudera:8020/user/hdfs/.Trash
22/07/28 00:33:54 INFO fs.TrashPolicyDefault: Created trash checkpoint: /user/hdfs/.Trash/220728003354
[cloudera@quickstart ~]$ █
```

## Creating new empty file in hdfs:

```
[cloudera@quickstart ~]$ sudo -u hdfs hadoop fs -touchz /user/19ce063/new-arpan.txt
[cloudera@quickstart ~]$ hadoop fs -ls /user/19ce063
Found 3 items
-rw-r--r-- 1 hdfs supergroup      18 2022-07-28 00:00 /user/19ce063/19ce063_arpan.txt
drwxr-xr-x - hdfs supergroup      0 2022-07-28 00:12 /user/19ce063/dummy
-rw-r--r-- 1 hdfs supergroup      0 2022-07-28 00:35 /user/19ce063/new-arpan.txt
```

To know the size of file:

```
[cloudera@quickstart ~]$ sudo -u hdfs hadoop fs -du /user/19ce063/
18 18 /user/19ce063/19ce063_arpan.txt
0 0 /user/19ce063/dummy
0 0 /user/19ce063/new-arpan.txt
[cloudera@quickstart ~]$ sudo -u hdfs hadoop fs -du /user/
18 18 /user/19ce063
18 18 /user/cloudera
18 18 /user/copying_in_anotherfolder.txt
18 18 /user/hdfs
0 0 /user/hduser
0 0 /user/history
0 0 /user/hive
0 0 /user/hue
0 0 /user/jenkins
858119738 858119738 /user/oozie
0 0 /user/root
0 0 /user/spark
```

```
[cloudera@quickstart ~]$ sudo -u hdfs hadoop fs -df /user/19ce063
Filesystem              Size      Used    Available  Use%
hdfs://quickstart.cloudera:8020 58479091712 869449728 46100221952    1%
```

To see statistics of folder:

```
[cloudera@quickstart ~]$ sudo -u hdfs hadoop fs -stat /user/19ce063
2022-07-28 07:35:45
```

To check health of file or folder:

```
[cloudera@quickstart ~]$ hadoop fsck /user/19ce063
DEPRECATED: Use of this script to execute hdfs command is deprecated.
Instead use the hdfs command for it.

Connecting to namenode via http://quickstart.cloudera:50070/fsck?ugi=cloudera&path=%2Fuser%2F19ce063
FSCK started by cloudera (auth:SIMPLE) from /127.0.0.1 for path /user/19ce063 at Thu Jul 28 01:07:32 PDT 2022
..Status: HEALTHY
Total size:      18 B
Total dirs:      2
Total files:      2
Total symlinks:    0
Total blocks (validated): 1 (avg. block size 18 B)
Minimally replicated blocks: 1 (100.0 %)
Over-replicated blocks: 0 (0.0 %)
Under-replicated blocks: 0 (0.0 %)
Mis-replicated blocks: 0 (0.0 %)
Default replication factor: 1
Average block replication: 1.0
Corrupt blocks: 0
Missing replicas: 0 (0.0 %)
Number of data-nodes: 1
Number of racks: 1
FSCK ended at Thu Jul 28 01:07:32 PDT 2022 in 2 milliseconds

The filesystem under path '/user/19ce063' is HEALTHY
```

## Change group with a new name (supergroup -> newgroup):

```
[cloudera@quickstart ~]$ sudo -u hdfs hadoop fs chgrp newgroup /user/19ce063/19ce063_arpan.txt
chgrp: Unknown command
Did you mean -chgrp? This command begins with a dash.
[cloudera@quickstart ~]$ sudo -u hdfs hadoop fs -chgrp newgroup /user/19ce063/19ce063_arpan.txt
[cloudera@quickstart ~]$ sudo -u hdfs hadoop fs -ls /user/19ce063
Found 3 items
-rw-r--r-- 1 hdfs newgroup      18 2022-07-28 00:00 /user/19ce063/19ce063_arpan.txt
drwxr-xr-x - hdfs supergroup    0 2022-07-28 00:12 /user/19ce063/dummy
-rw-r--r-- 1 hdfs supergroup    0 2022-07-28 00:35 /user/19ce063/new-arpan.txt
```

To merge content of two file into one file:





**Aim:**

**Practical 3**

**To understand the overall programming architecture using Map Reduce API.(word count using Map-Reduce)**

**Understand flow of MR,**

- **Create mapper and reducer file**
- **Set MR execution from CMD over HDFS**

**mapper.py**

```
#!/usr/bin/python
import sys
for line in sys.stdin:
    line = line.strip()
    words = line.split()
    for word in words:
        print('%s\t%s' % (word, 1))
```

**reducer.py**

```
#!/usr/bin/python

from operator import itemgetter
import sys

current_word = None
current_count = 0
word = None

for line in sys.stdin:
    line =
```

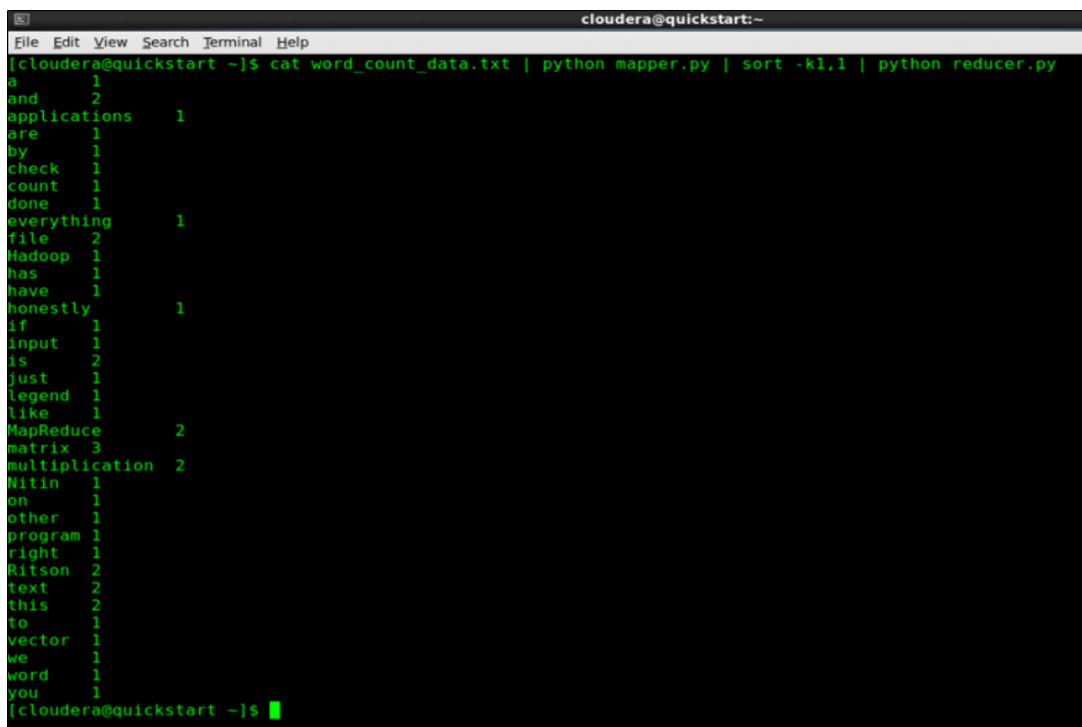
```

line.strip()    word, count =
line.split('\t', 1)    try:
    count = int(count)
except ValueError:
    continue

    if current_word == word:
current_count += count
else:    if current_word:
    print('%s\t%s' % (current_word, current_count))
current_count = count    current_word = word if
current_word == word:
    print('%s\t%s' % (current_word, current_count))

```

## ➔ Output



A terminal window titled 'cloudera@quickstart:~' showing the execution of a word count program. The command executed is `cat word_count_data.txt | python mapper.py | sort -k1,1 | python reducer.py`. The output displays words and their counts, sorted by word. The words and counts are: a (1), and (2), applications (1), are (1), by (1), check (1), count (1), done (1), everything (1), file (2), Hadoop (1), has (1), have (1), honestly (1), if (1), input (1), is (2), just (1), legend (1), like (1), MapReduce (2), matrix (3), multiplication (2), Nitin (1), on (1), other (1), program (1), right (1), Ritson (2), text (2), this (2), to (1), vector (1), we (1), word (1), you (1).

```

cloudera@quickstart:~
File Edit View Search Terminal Help
[cloudera@quickstart ~]$ cat word_count_data.txt | python mapper.py | sort -k1,1 | python reducer.py
a 1
and 2
applications 1
are 1
by 1
check 1
count 1
done 1
everything 1
file 2
Hadoop 1
has 1
have 1
honestly 1
if 1
input 1
is 2
just 1
legend 1
like 1
MapReduce 2
matrix 3
multiplication 2
Nitin 1
on 1
other 1
program 1
right 1
Ritson 2
text 2
this 2
to 1
vector 1
we 1
word 1
you 1
[cloudera@quickstart ~]$

```

**Aim:****Practical 4****Configure Hive and perform SQL like operation Steps:****1) Open HDFS in Cloudera and create csv file.**

```

[cloudera@quickstart ~]$ ls
cloudera-manager  Downloads  kerberos  parcels  Videos
cm_api.py         eclipse    lib        Pictures workspace
Desktop          enterprise-deployment.json  Music     Public
Documents        express-deployment.json    myNewFile Templates
[cloudera@quickstart ~]$ cd Documents
[cloudera@quickstart Documents]$ ls
cloudera-manager.html  Customer.csv  Employee.csv  Order.csv
css                   employee.csv~  img           Order.csv~
[cloudera@quickstart Documents]$ cat Employee.csv
Id,Name,Dept,Yoj,Salary
1,Rose ,IT ,2012,26000
2,Sam,Sales,2012,22000
3,Mike,HR,2013,30000
4,Nick,SC ,2013,20000
[cloudera@quickstart Documents]$

```

**2) Open HIVE in command prompt****3) Create Database**

```
hive> CREATE DATABASE office;
```

**4) Use database**

```
hive> USE office;
```

**5) Create Table**

```

office
Time taken: 0.02 seconds, Fetched: 2 row(s)
hive> create table employee
> (Id INT, Name STRING, Dept STRING, Yoj INT, salary INT)
> row format delimited fields terminated by ','
> tblproperties ("skip.header.line.count"="1");

```

**6) Insert Data from the created csv file into created table.**

```

Time taken: 0.271 seconds
hive> show tables;
OK
employee
Time taken: 0.056 seconds, Fetched: 1 row(s)
hive> describe employee;
OK
id                int
name              string
dept              string
yobj              int
salary            int
Time taken: 0.201 seconds, Fetched: 5 row(s)
hive> select * from employee;
OK
Time taken: 0.621 seconds
hive> LOAD DATA LOCAL INPATH
> '/home/cloudera/Documents/Employee.csv'
> INTO TABLE employee;
Loading data to table office.employee
Table office.employee stats: [numFiles=1, totalSize=116]
OK
Time taken: 1.147 seconds
hive>

```

## 7) Show data and perform different operation.

### Select, Show table, alter table, drop table.

```

hive> select * from office.employee WHERE Salary>25000;
OK
1      Rose    IT      2012    26000
3      Mike    HR      2013    30000
Time taken: 0.296 seconds, Fetched: 2 row(s)
hive> alter table office.employee RENAME TO office.employees;
OK
Time taken: 0.496 seconds
hive> show tables;
OK
employees
Time taken: 0.024 seconds, Fetched: 1 row(s)
hive> drop table employees;
OK
Time taken: 0.363 seconds
hive> show tables;
OK
Time taken: 0.017 seconds
hive>

```

## Practical 5

### Configure Spark and perform Action and Transformation on RDD

#### Program:

## Aim:

```
val bot = List((1,"S"),(2,"Sagar")) val rdd =
sc.parallelize(bot) rdd.foreach(println)
println("Transformation and Action:")
println("With map") val rdd2=rdd.map(f=>
(f,1)) rdd2.foreach(println) println("With map
Key")
val rdd3=rdd.map(f=> (f,1)).sortByKey()
rdd3.foreach(println) val c = rdd2.count()
println("Count:")
println(c)
```

## Output:

```
Administrator: Command Prompt - spark-shell -i rdd2.scala
asses where applicable
Using Spark's default log4j profile: org/apache/spark/log4j-defaults.properties
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
Spark context Web UI available at http://6168-05:4040
Spark context available as 'sc' (master = local[*], app id = local-1665480759494).
Spark session available as 'spark'.
(1,S)e 0:> (0 + 12) / 12]
(2,Sagar)
Transformation and Action:
With map
((2,Sagar),1)
((1,S),1)
With map Key
((2,Sagar),1)
((1,S),1)
Count:
2
Welcome to
  ____  _
 / ___|| | | |
| |___| |_| |
 \___|_||_|_|_|
version 3.1.3

Using Scala version 2.12.10 (Java HotSpot(TM) 64-Bit Server VM, Java 1.8.0_211)
Type in expressions to have them evaluated.
Type :help for more information.

scala>
```

## Practical 6

**Aim: Perform SparkSQL on Netflix\_title dataset(Given) implementation:**



```
movies = spark.read.format("csv") \
.option("header", "true") \
.option("inferSchema", "true") \
.load("../input/netflix-files-tests/netflix_titles.csv")
```

```
movies.printSchema()
```

```
root
 |-- show_id: string (nullable = true)
 |-- type: string (nullable = true)
 |-- title: string (nullable = true)
 |-- director: string (nullable = true)
 |-- cast: string (nullable = true)
 |-- country: string (nullable = true)
 |-- date_added: string (nullable = true)
 |-- release_year: string (nullable = true)
 |-- rating: string (nullable = true)
 |-- duration: string (nullable = true)
 |-- listed_in: string (nullable = true)
 |-- description: string (nullable = true)
```

```
movies.show(3)
```

```
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+
|show_id|  type|title|      director|      cast|  country|      date_added|release_year|rating| duration|
listed_in|      description|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+
|   s1|TV Show|  3%|      null|João Miguel, Bian...|  Brazil|  August 14, 2020|      2020|  TV-MA|4 Seasons|Intern
ational TV ...|In a future where...|
|   s2|  Movie| 7:19|Jorge Michel Grau|Demián Bichir, Hé...|  Mexico|December 23, 2016|      2016|  TV-MA|  93 min|Drama
s, Internati...|After a devastati...|
|   s3|  Movie|23:59|  Gilbert Chan|Tedd Chan, Stella...|Singapore|December 20, 2018|      2011|    R|  78 min|Horror
Movies, In...|When an army recr...|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+
only showing top 3 rows
```

```
df = movies.select('title', 'release_year', 'country', 'rating') df.show(4)
```

```

+-----+-----+-----+-----+
|title|release_year|country|rating|
+-----+-----+-----+-----+
| 3%|2020|Brazil|TV-MA|
| 7:19|2016|Mexico|TV-MA|
|23:59|2011|Singapore|R|
| 9|2009|United States|PG-13|
+-----+-----+-----+-----+

```

only showing top 4 rows

```
df.printSchema()
```

```

root
 |-- title: string (nullable = true)
 |-- release_year: string (nullable = true)
 |-- country: string (nullable = true)
 |-- rating: string (nullable = true)

```

```
df2 = df.withColumn('year', df['release_year'].cast('int')).drop('release_year')
df2.printSchema()
```

```

root
 |-- title: string (nullable = true)
 |-- country: string (nullable = true)
 |-- rating: string (nullable = true)
 |-- year: integer (nullable = true)

```

```
df2.filter('year > 2015').show(5)
```

```

+-----+-----+-----+-----+
|title|country|rating|year|
+-----+-----+-----+-----+
| 3%|Brazil|TV-MA|2020|
| 7:19|Mexico|TV-MA|2016|
| 46|Turkey|TV-MA|2016|
| 122|Egypt|TV-MA|2019|
| 706|India|TV-14|2019|
+-----+-----+-----+-----+
only showing top 5 rows

```

```
# min e max
from pyspark.sql.functions import max, min
```

```
df2.select(max('year')).show(3)
```

```
+-----+
|max(year)|
+-----+
|      2021|
+-----+
```

```
df2.filter('year == 2021 and country is not NULL').show(10)
```

```
+-----+-----+-----+-----+
|          title|          country|rating|year|
+-----+-----+-----+-----+
|   Carmen Sandiego|   United States|TV-Y7|2021|
|    Charming|Canada, United St...|TV-Y7|2021|
|    Cobra Kai|   United States|TV-14|2021|
|Crack: Cocaine, C...|   United States|TV-MA|2021|
|   Disenchantment|   United States|TV-14|2021|
| Dream Home Makeover|   United States|TV-G|2021|
|Headspace Guide t...|   United States|TV-G|2021|
|          Hilda|United Kingdom, C...|TV-Y7|2021|
|History of Swear ...|   United States|TV-MA|2021|
|Inside the World'...|   United Kingdom|TV-MA|2021|
+-----+-----+-----+-----+
only showing top 10 rows
```

## Practical 7

**Aim: Perform Spark streaming with word count.**

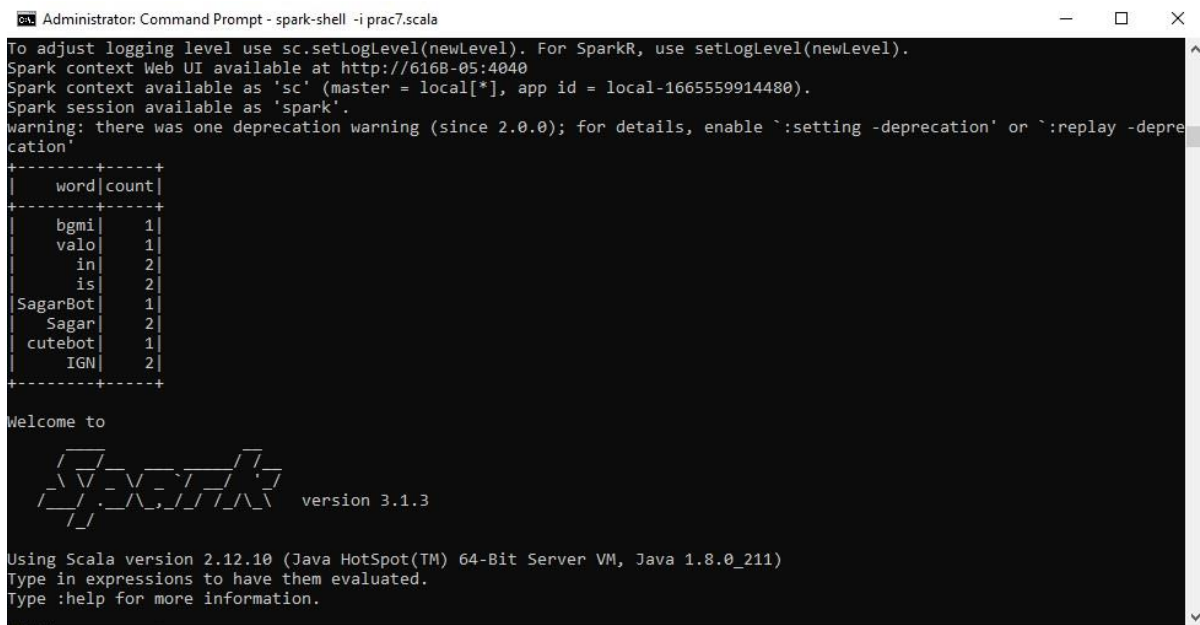
**Program:**

```
val linesDF = sc.textFile("mytext.txt").toDF("line")
val wordsDF = linesDF.explode("line", "word")((line: String) => line.split(" "))
val wordCountDF = wordsDF.groupBy("word").count() wordCountDF.show()
```

## Input:

```
Sagar IGN is cutebot in bgmi
Sagar IGN is SagarBot in valo
```

## Output:

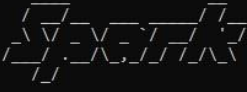


Administrator: Command Prompt - spark-shell -i prac7.scala

To adjust logging level use `sc.setLogLevel(newLevel)`. For SparkR, use `setLogLevel(newLevel)`.  
 Spark context Web UI available at <http://6168-05:4040>  
 Spark context available as 'sc' (master = local[\*], app id = local-1665559914480).  
 Spark session available as 'spark'.  
 warning: there was one deprecation warning (since 2.0.0); for details, enable `:setting -deprecation` or `:replay -deprecation`

word	count
bgmi	1
valo	1
in	2
is	2
SagarBot	1
Sagar	2
cutebot	1
IGN	2

Welcome to

 version 3.1.3

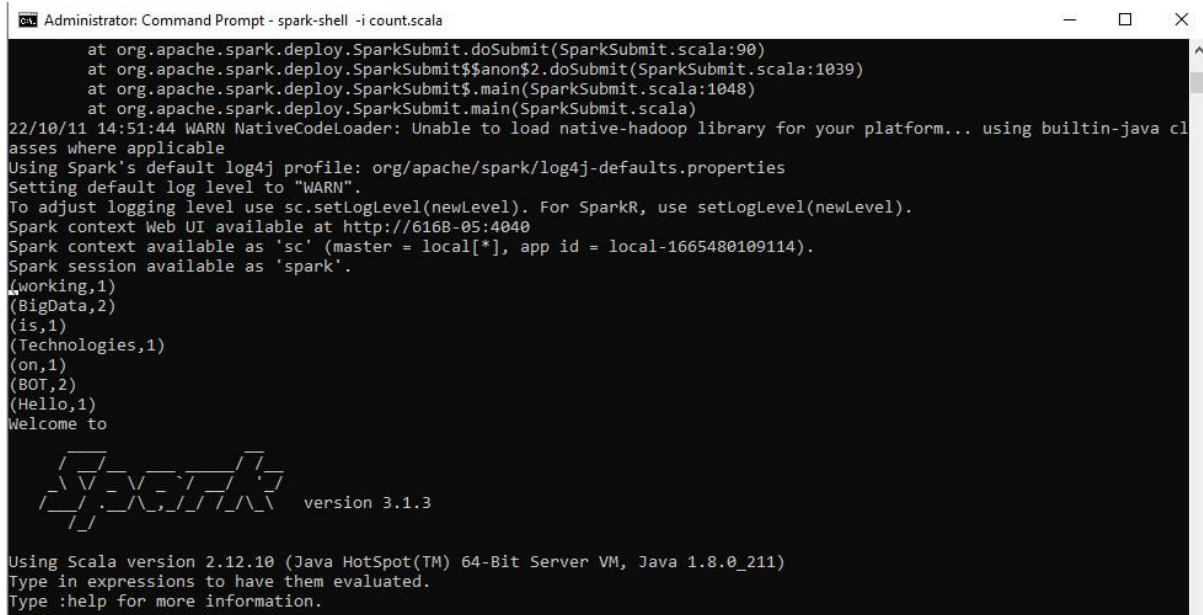
Using Scala version 2.12.10 (Java HotSpot(TM) 64-Bit Server VM, Java 1.8.0\_211)  
 Type in expressions to have them evaluated.  
 Type :help for more information.

## Practical 8

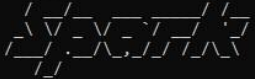
**Aim: Perform word count operation on it using SCALA Program:**

```
val list = List("BOT is working on BigData Technologies","Hello BOT","BigData") val
words = list.flatMap(line => line.split(" ")) val keyData = words.map(word =>
(word,1)) val groupedData = keyData.groupBy(_._1) val result =
groupedData.mapValues(list=>{ list.map(_._2).sum
})
result.foreach(println)
```

**Output:**



```
Administrator: Command Prompt - spark-shell -i count.scala
at org.apache.spark.deploy.SparkSubmit.doSubmit(SparkSubmit.scala:90)
at org.apache.spark.deploy.SparkSubmit$$anon$2.doSubmit(SparkSubmit.scala:1039)
at org.apache.spark.deploy.SparkSubmit$.main(SparkSubmit.scala:1048)
at org.apache.spark.deploy.SparkSubmit.main(SparkSubmit.scala)
22/10/11 14:51:44 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java cl
asses where applicable
Using Spark's default log4j profile: org/apache/spark/log4j-defaults.properties
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
Spark context Web UI available at http://6168-05:4040
Spark context available as 'sc' (master = local[*], app id = local-1665480109114).
Spark session available as 'spark'.
(working,1)
(BigData,2)
(is,1)
(Technologies,1)
(on,1)
(BOT,2)
(Hello,1)
Welcome to

 version 3.1.3

Using Scala version 2.12.10 (Java HotSpot(TM) 64-Bit Server VM, Java 1.8.0_211)
Type in expressions to have them evaluated.
Type :help for more information.
```

## Practical – 9

➔ **AIM** : Configure Neo4j and create node(add –remove property) & relationship in it.

---

➔ **Steps to Configure Neo4J** :

1.Download Neo4J Community Edition from below link:

<https://neo4j.com/download-center/>

2.Extract the zip file and keep the extracted folder in D:\drive (for e.g.)

3.Open command prompt and navigate to D:\neo4j

4.Run below commands to start neo4j service:bin\neo4j  
installservicebin\neo4j start

5.Open the URL <http://localhost:7474/browser/> in your web browser.

6.Initial username: neo4j and password: neo4j

7.Once submitted prompt for change password will be shown wherein you can change database password

8.Homepage will be shown in which top most text box starting with \$ is for writing queries.

➔ **Code sample** :

**(for implementation)**

CREATE (n: Person {name:"Brad"}) RETURN n

CREATE (n: Person {name:"Alice"}) RETURN n

CREATE (n: Person {name:"Mike"}) RETURN n

CREATE (n: Person {name:"Jill"}) RETURN n

CREATE (n: Person {name:"Hazel"}) RETURN n

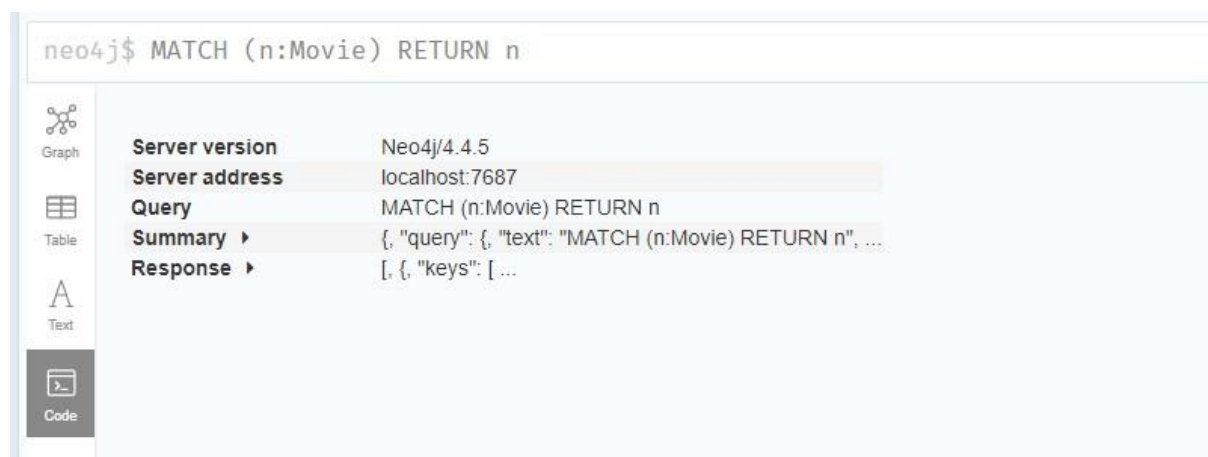
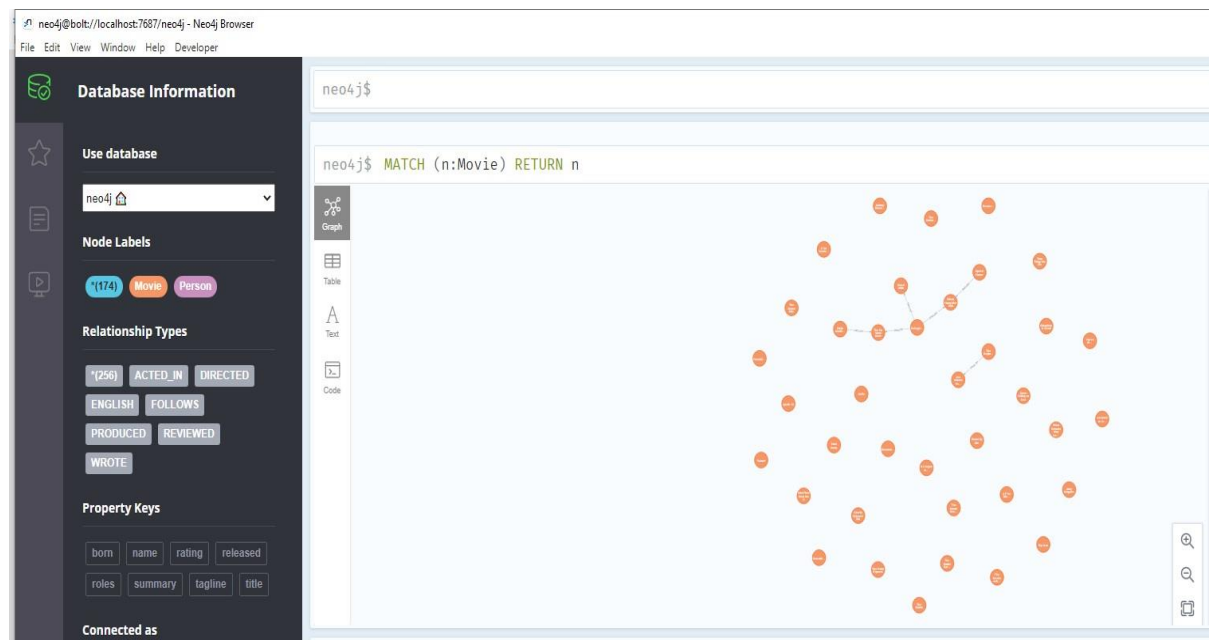
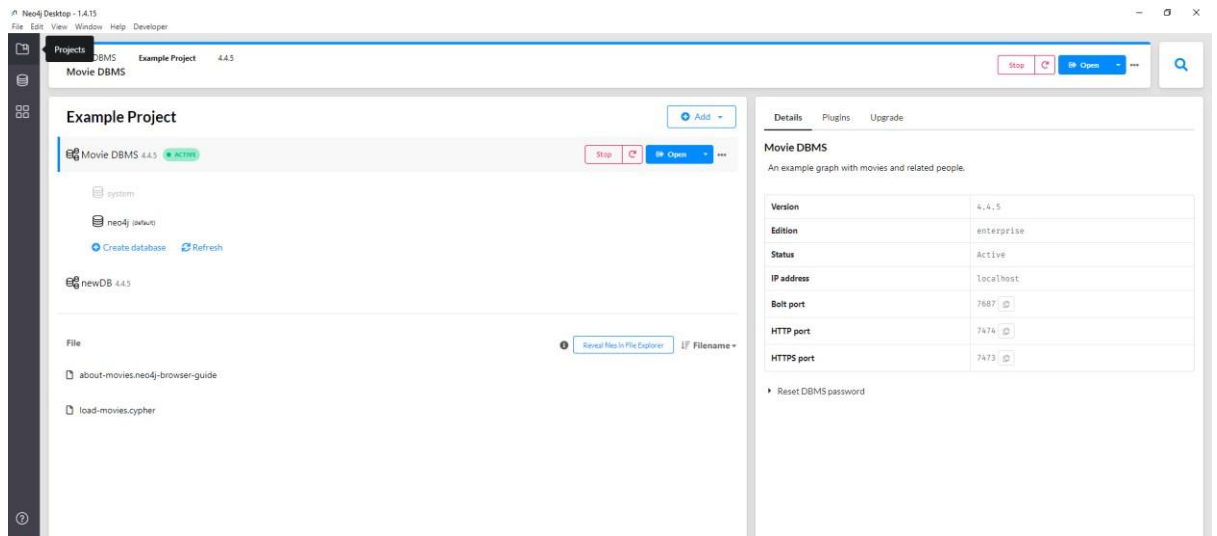
CREATE (n: Movie {title:"Unforgiven "}) RETURN n

CREATE (n: Movie {title:"The Da Vinci Code"}) RETURN n

CREATE (n: Movie {title:"When Harry Met Sally"}) RETURN n ➔

**Output Screenshots** :





## Practical – 10

➔ **AIM** : Import files from Neo4j and complete relational graph from it.

➔ **Code sample:**  
(for implementation)

1.To add a relationship to nodes:

```
MATCH (a:Person{name:"Brad"}),(b:Person{name:"Alice"}) MERGE (a) -[r:FRIENDS]-> (b)
```

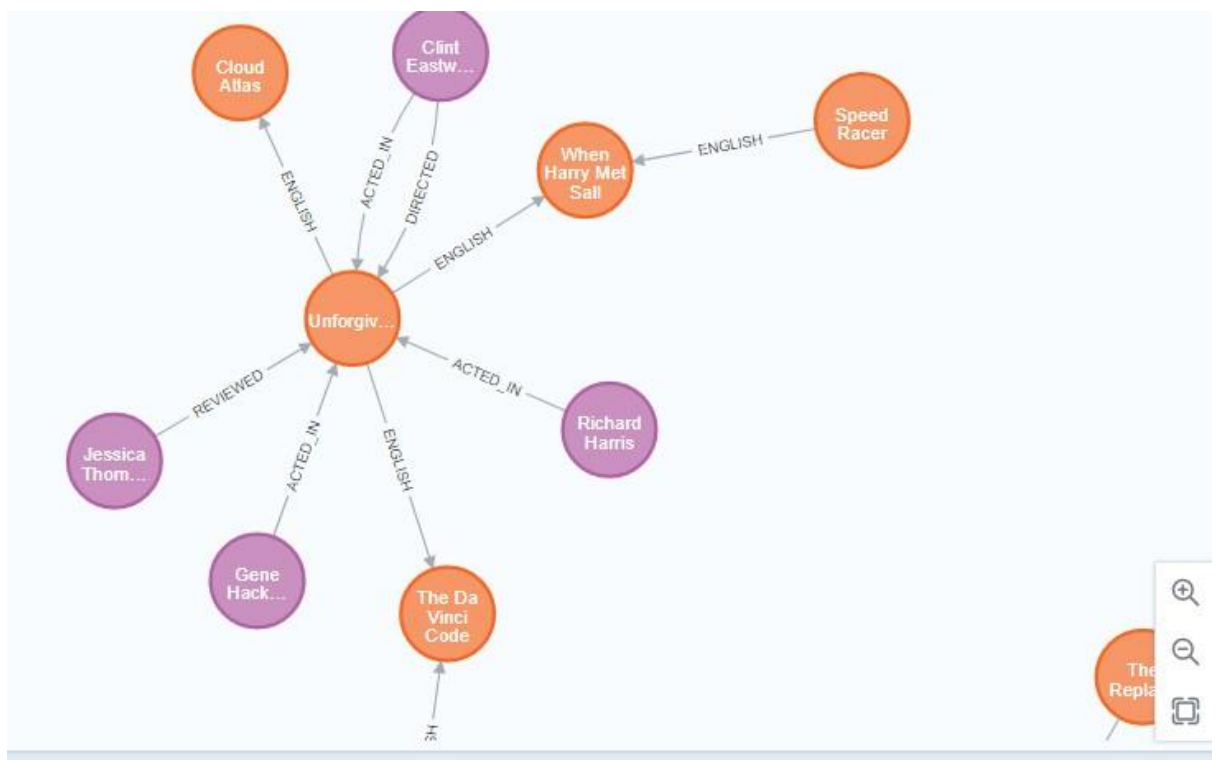
2.To add properties in relationship:

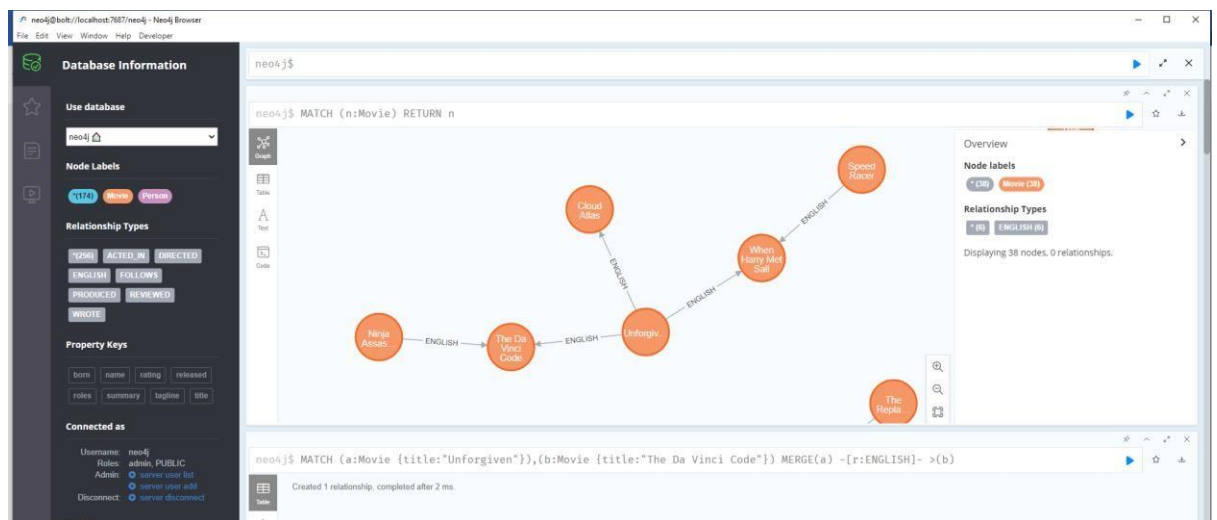
```
MATCH (a: Person {name:"Brad"}),(b: Person {name:"MIKE"}) MERGE (a) -[r: FRIENDS {since:"1998"}]-> (b)
```

3.Adding relationship between two different labels:

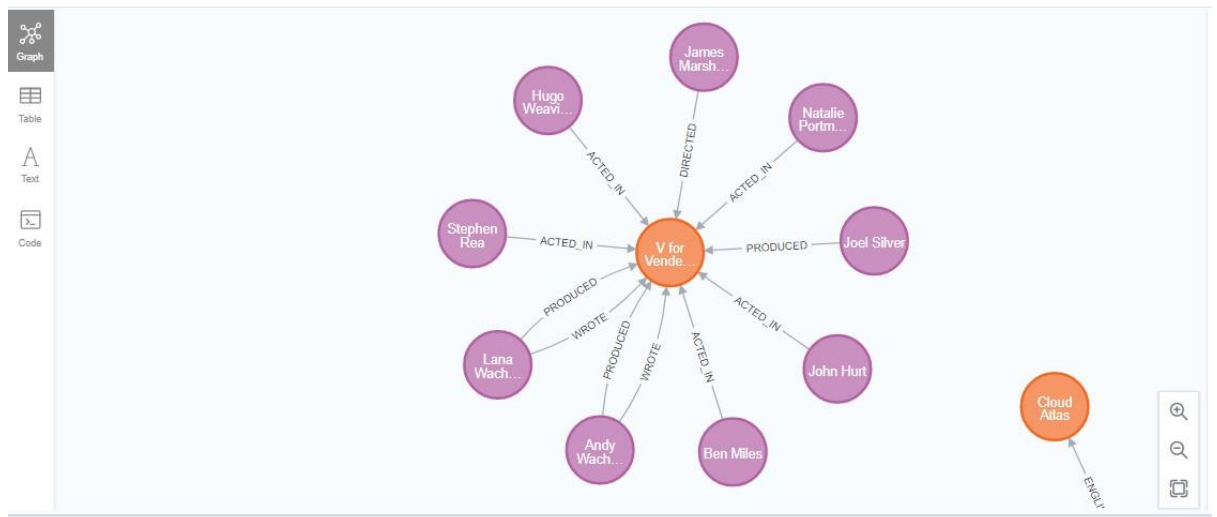
```
MATCH (a:Person{name:"Jill"}),(b:Movie{name:"Avengers"}) MERGE (a) -[r:FAVOURITE]-> (b)
```

➔ **Implementation Screenshots :**





neo4j\$ MATCH (n:Movie) RETURN n



A	B
C	D

E

F

**\*\*Description\*\***

Table 1: Name of table

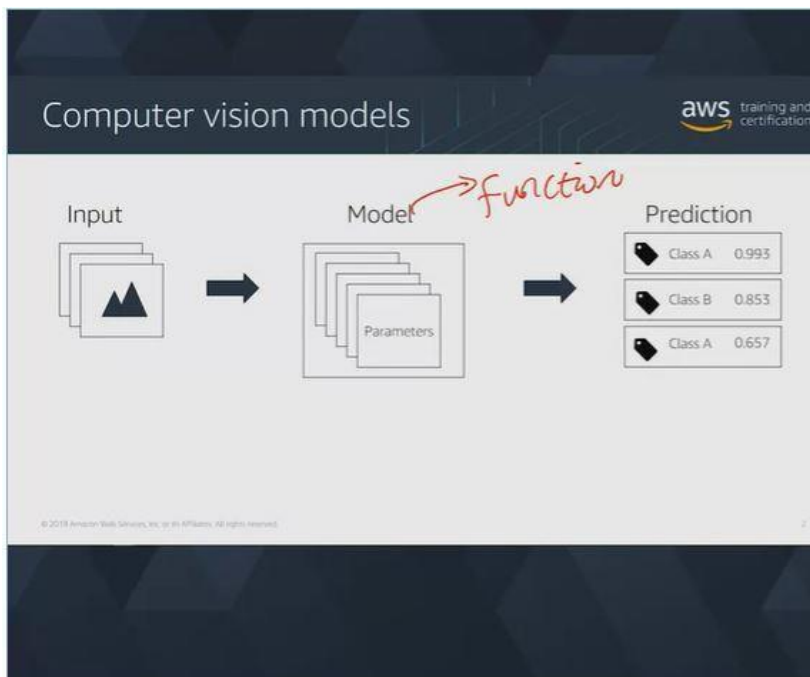


Figure 1: Name of image/figure/diagram